

# CS240 Assignment 4

## Fall 2013, Prof. Hany Ramadan

Emaad Ahmed Manzoor (129110)

November 28, 2013

The following describes the per-process and global state, and associated methods that manipulate these pieces of state.

**Global State.** Each `struct semaphore` is present in a `MAXSEMS` size array in the kernel, and maintains the following fields:

- `uint name`: Initially 0, this records the name provided to the `sem_get` call.
- `int wakeups`: Initially 0, this records the value of the semaphore, set by a `sem_get` call.
- `uint used`: Initially 0, this records the number of processes using this semaphore. It is the total number of `sem_gets` minus the total number of `sem_deletes` on this semaphore.

**Per-process State.** An integer array `proc->sems` of `MAXSEMS` size is maintained, with each element either set to 0 or 1. This is to prevent processes from calling semaphore functions on a handle without first calling `sem_get` to obtain the handle.

There is also a lock protecting the semaphore array in the kernel called `semlock`.

**Functions.** There are four functions that manipulate this state. All functions check if the `handle` is in the correct range, if the semaphore corresponding to the `handle` is `used`, and (except `sem_get`) if the calling process has its `proc->sems[handle]` flag set to 1.

`sem_get_proc(name, value, proc)` finds a new handle or returns an existing handle for a semaphore. If the `proc->sems[handle]` for the provided `proc` structure is 0, it is set to 1 and the `used` count for the semaphore is incremented.

`sem_delete_proc(handle, proc)` decrements the `used` count for the semaphore and sets `proc->sems[handle]` to 0 for the provided `proc` structure. It then wakes up all processes waiting on this semaphore. Note that this ensures a process can call `sem_delete` only once on a semaphore obtaining via `sem_get`. Once all such processes have `exited` or called `sem_delete`, the `used` count becomes 0 and this semaphore can be used by another `sem_get` call.

`sem_wait(handle)` runs a loop checking the number of `wakeups` of the semaphore. If the number of `wakeups` is 0, it `sleeps` with the address of the `handleth` element in the kernel semaphore array as the wait channel, and `semlock` as the lock. Once more than 0 `wakeups` are observed, `wakeups` is decremented and the function breaks out of the loop.

`sem_signal(handle)` increments the number of `wakeups` for the semaphore and calls `wakeup` with the address of the `handleth` element in the kernel semaphore array as the wait channel.

**Fork.** On forking, `sem_get` is simulated on the child for every semaphore present in the parent's `proc->sems` array. This is done by calling `sem_get_proc` for every parent semaphore and providing the child's `proc` structure with the parent's semaphore's name.

**Exit.** On exit, `sem_delete` is simulated on every semaphore present in the parent's `proc->sems` array. This is done by calling `sem_delete_proc` and providing the parent's `proc` structure with the handles present in its `proc->sems` array.

**Limits.** `MAXSEMS` is the maximum number of semaphores in the system, defined in `params.h`.

**Tests.** Semaphores were implemented on top of the shared memory implemented in the previous assignment. The provided producer-consumer user program was modified to use shared memory.

### Files Modified or Created.

Makefile	Entries for new source files.
defs.h syscall.c syscall.h sysproc.c user.h usys.S	Stub code and function prototypes for the system calls.
semaphore.c semaphore.h proc.c proc.h semaphoretest.c param.h	Semaphores implementation, parameters, process semaphore array and modified fork and exit functions.